

Portfolio 2 — Investment Data Platform

This document describes the scope, objectives, and delivery roadmap for the Investment Data Platform portfolio project, built using Microsoft Fabric.

Project Objective

Design and implement an end-to-end data engineering platform for financial market data, focused on ingestion, transformation, modeling, and analytics-ready consumption.

Core Principles

- 1 End-to-end data engineering focus (not a simple tracker)
- 2 Production-oriented architecture using Microsoft Fabric
- 3 Clear separation of raw, curated, and analytics layers
- 4 Business-oriented financial use case

Architecture Overview

The project follows a Lakehouse architecture with Medallion layers: Bronze for raw ingestion, Silver for cleaned and standardized data, and Gold for analytics-ready datasets consumed by Power BI.

Technology Stack

- 1 Microsoft Fabric (Lakehouse, Data Pipelines, Notebooks)
- 2 PySpark and SQL for transformations
- 3 Public financial market APIs
- 4 Power BI for analytics consumption

Roadmap

- 1 **Week 1 — Foundation:** Data source selection, repository setup, Lakehouse creation, Bronze ingestion pipeline.
- 2 **Week 2 — Processing:** Silver layer transformations, data normalization, historical handling.
- 3 **Week 3 — Analytics:** Gold layer modeling, financial metrics calculation, BI-ready structures.
- 4 **Week 4 — Delivery:** Power BI dashboard, documentation, final review.

Expected Outcomes

- 1 A complete and documented data engineering project
- 2 Clear demonstration of Microsoft Fabric capabilities
- 3 Analytics-ready financial datasets
- 4 Interview-ready portfolio artifact